



Why Cosmos DB?

What traditional databases were lacking?



Challenges with globally distributed Databases

- Long time
- Lot of effort
- Need own infrastructure
 - Teams
 - Data centers etc.





How Cosmos DB evolved?

<u>In 2010</u>

Microsoft realized they need to build something very different to handle global distribution need.

<u>In 2015</u>

Microsoft released Azure
Document DB

Supports SQL queries over JSON documents

<u>In 2017</u>

Rebranded Document DB as Cosmos DB

Globally distributed and scalable database



Why Cosmos DB?

FULLY MANAGED

- Database as a service (DaaS)
 - Serverless architecture
 - No operational overhead
- No schema or Index management

GLOBALLY DISTRIBUTED

• Turnkey global distribution

Azure Cosmos DB's support f

 Azure Cosmos DB's support for consistency levels like strong, eventual, consistent prefix, session, and bounded-staleness.

SCALABLE

• Unlimited scale for both storage and throughput.

CONSISTENCY CHOICES

MULTIMODEL & MULTI-LANGUAGE

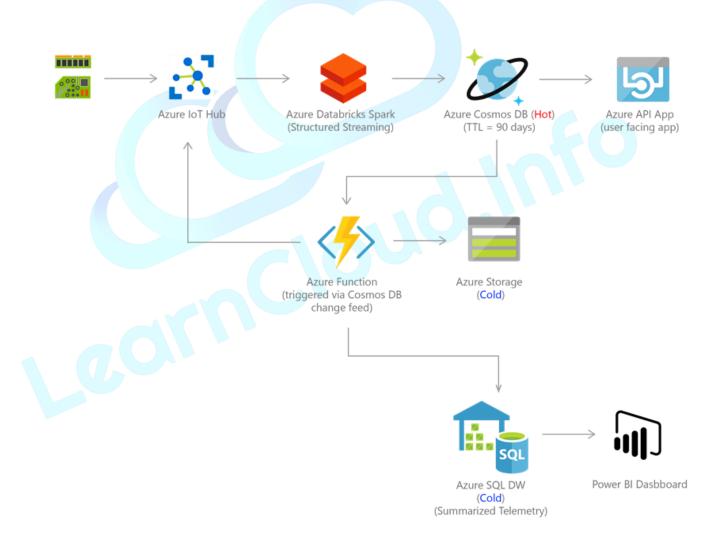
- Supports Jason documents, table graph and columnar data models
 - Java, .NET, Python, Node.js, JavaScript, etc.

HIGHLY AVAILABLE, RELIABLE & SECURE

- Always on
- 99.999% SLA
- < 10ms latency

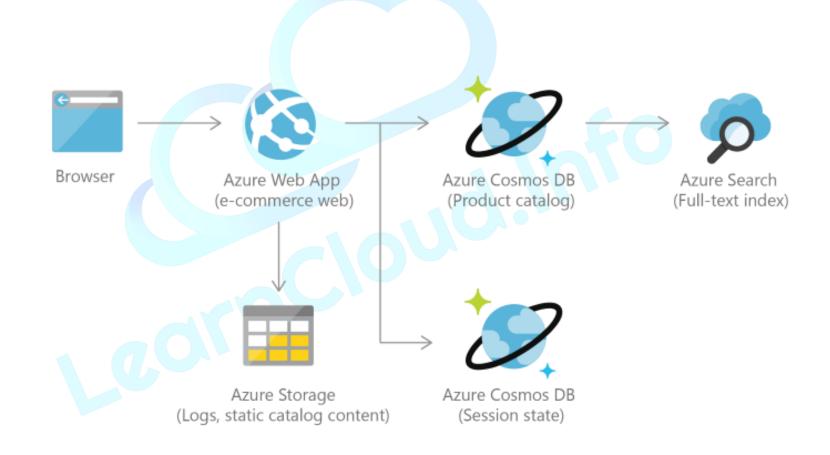


Use case - IOT



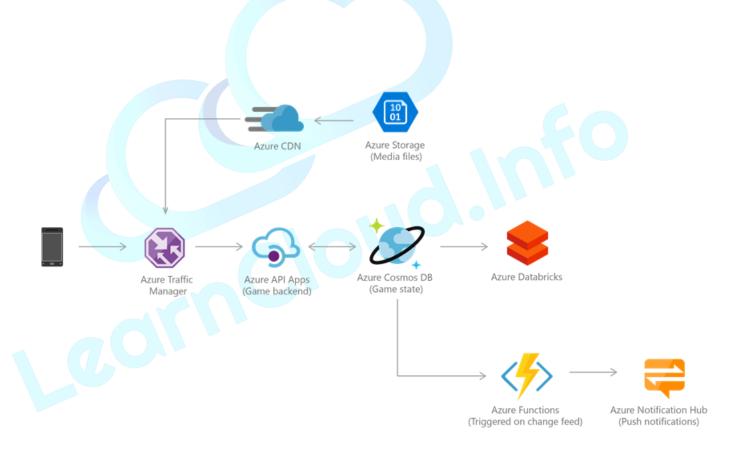


Use case - Retail and Marketing



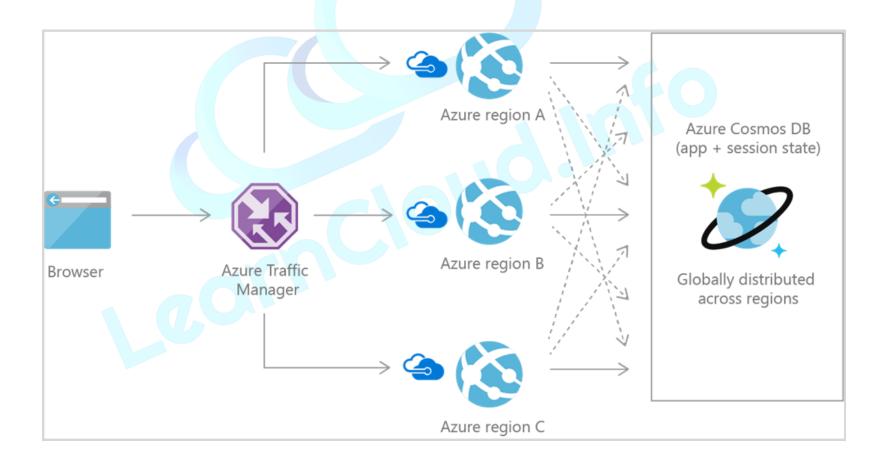


Use case - Gaming

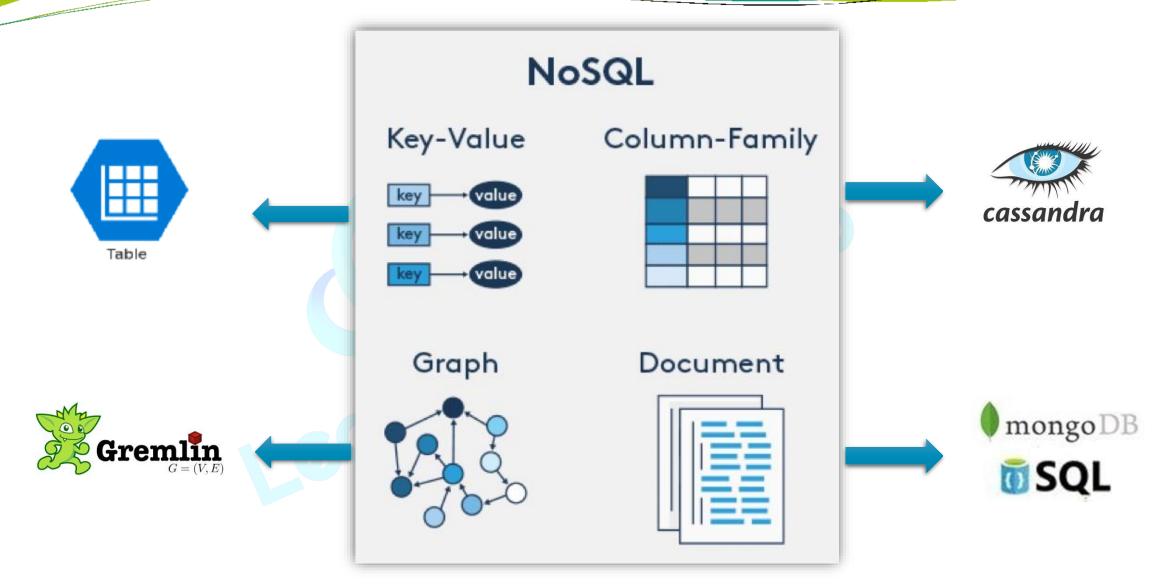




Use case - Web and mobile









SQL API vs MongoDB API

SQL(CORE) API

JSON Documents

Microsoft original Document DB platform Supports server side programming model

You can use SQL like language to query JSON documents.

MongoDB API

BSON Documents

Implement Wire protocol

Fully compatible with Mongo DB application code

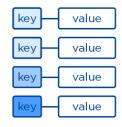
Migrate existing Cosmos DB without much change of logic

Use SQL(CORE) API for new development



Cosmos DB Table API

Key-Value





- Key-Value store
- Premium offering for Azure Table Storage
- Existing Table Storage customers will migrate to Cosmos DB Table API
- Row value can be simple like number or string
- Row cannot store object



Cosmos DB Cassandra API

Wide-column



cassandra

- Wide column No SQL Database
- Name and format of column can vary from row to row.
- Simple migrate your Cassandra application to Cosmos Cassandra API and change connection string.
- Interact
 - Cassandra based tools
 - Data Explorer
 - Programmatically, using SDK (CassandraCSharpdriver)



Cosmos DB Gremlin API

Graph

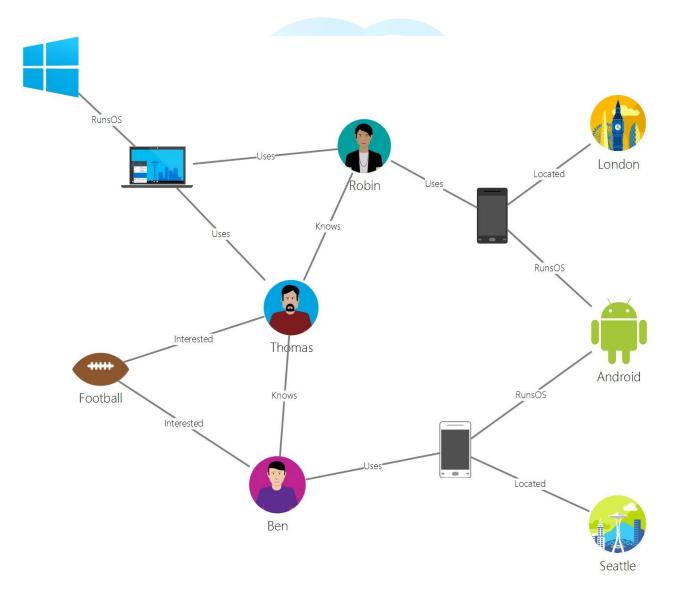


- Graph Data Model
- Real world data connected with each other
- Graph database can persist relationships in the storage layer





Graph Model





Cosmos DB Gremlin API

Graph





- Graph Data Model
- Real world data connected with each other
- Graph database can persist relationships in the storage layer
- Use cases
 - Social networks
 - Recommendation engines
 - Geospatial
 - Internet of things
- Migrate existing apps to Cosmos DB Gremlin API
- Graph traverse a language



Analyze the decision criteria

	Core (SQL)	MongoDB	Cassandra	Azure Table	Gremlin
New projects being created from scratch	✓				
Existing MongoDB, Cassandra, Azure Table, or Gremlin data		✓	✓	√	✓
Analysis of the relationships between data					√
All other scenarios	✓				



Azure Table storage vs Cosmos DB Table API

Azure Table Storage

- **➢** Geo replication is restricted
 - Only 1 additional pair region
- > Support for primary key lookups only
- Price optimized for cold storage
- Lower performance
 - Throughput is capped
 - Latency is higher
- No consistency options

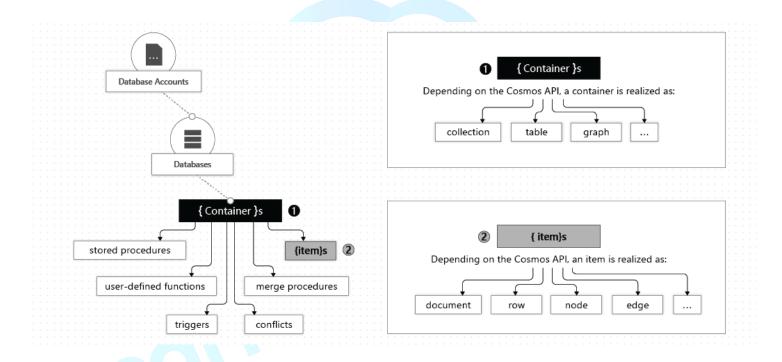


ငြစ္နန္ကာဝွန DB Table API

- Geo replication across your choice of any number of regions
- Secondary index support for lookups across multiple dimensions
- Better performance
 - Unlimited and predictable throughput
 - latency is lower
- 5 consistency options



Database Containers and Items



Azure Cosmos entity	SQL API	Cassandra API	MongoDB API	Gremlin API	Table API
Azure Cosmos database	Database	Keyspace	Database	Database	NA
Azure Cosmos container	Container	Table	Collection	Graph	Table
Azure Cosmos item	Document	Row	Document	Node or edge	Item

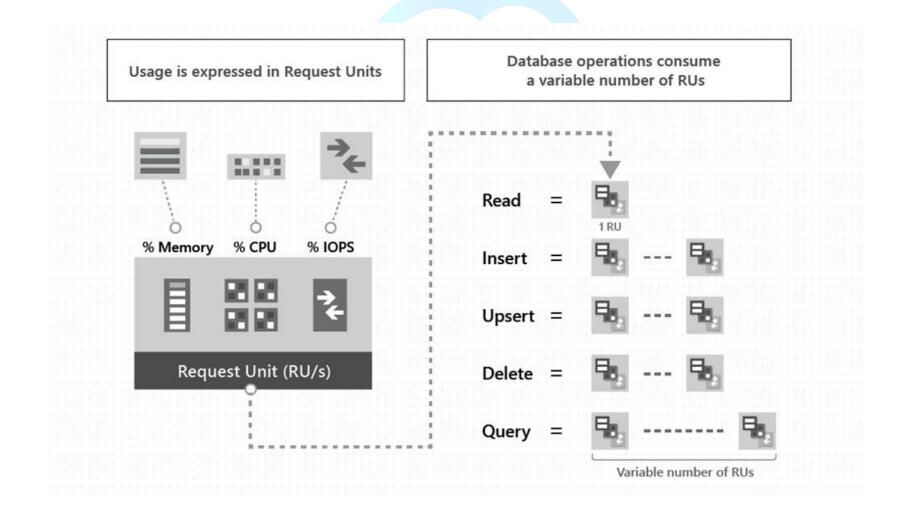


Measuring Performance

Latency - How fast is the response for a given request?
 How many request can be served within a specific period of time?

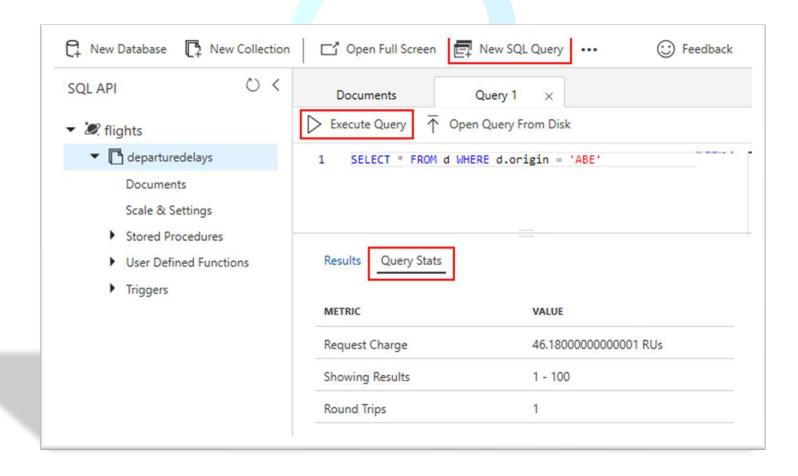


Introducing Request Units





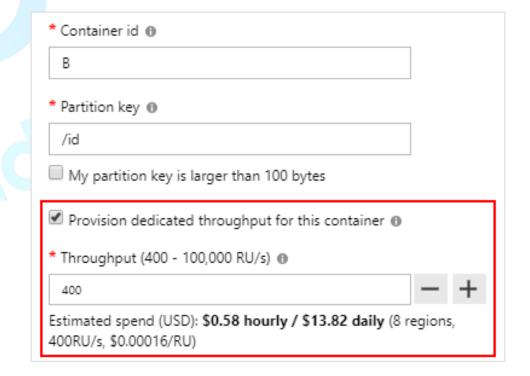
Introducing Request Units





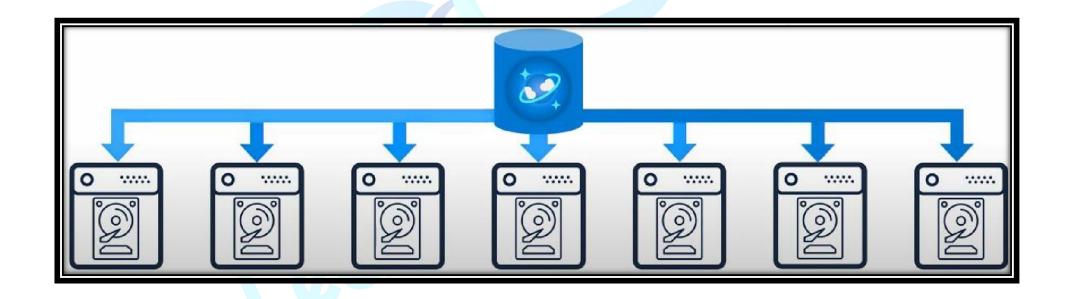
Reserving requests units

- Provision Request units per second (RU/s)
 - How many request units (not requests) per second are available to your application
- Exceeding reserved throughput limits
 - Requests are "throttled" (HTTP 429)





Horizontally Scalable

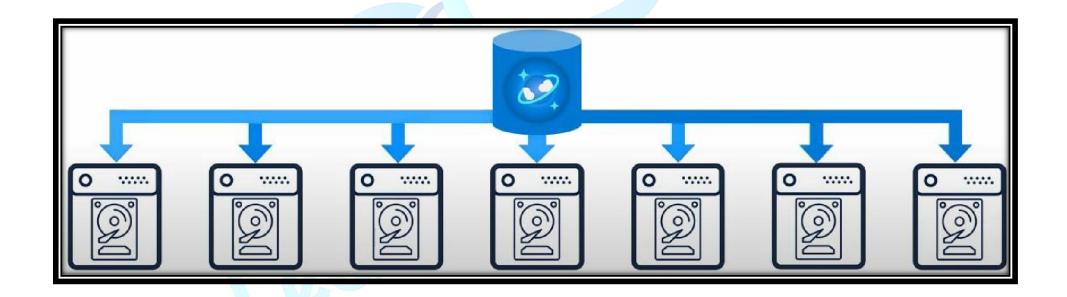


Unlimited Storage

Unlimited Throughput



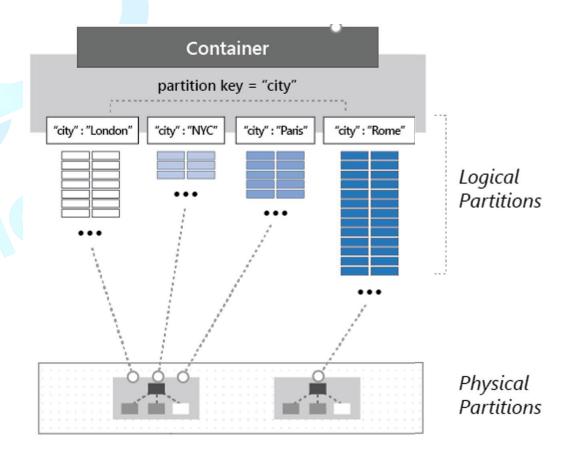
Partitioning





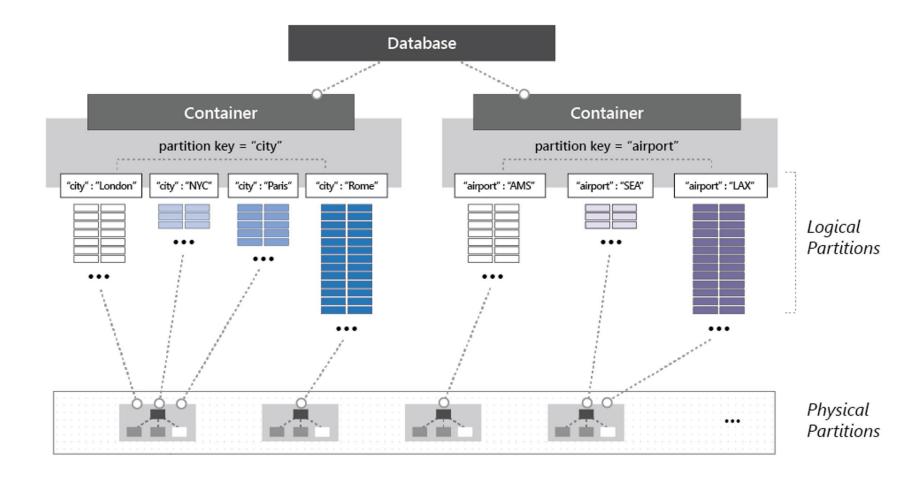
Partitioning

- Partitioning: the items in a container are divided into distinct subsets called logical partitions.
- Partition key is the value by which Azure organizes your data into logical divisions.
- Logical partitions are formed based on the value of a partition key that is associated with each item in a container.
- Physical partitions: Internally, one or more logical partitions are mapped to a single physical partition.



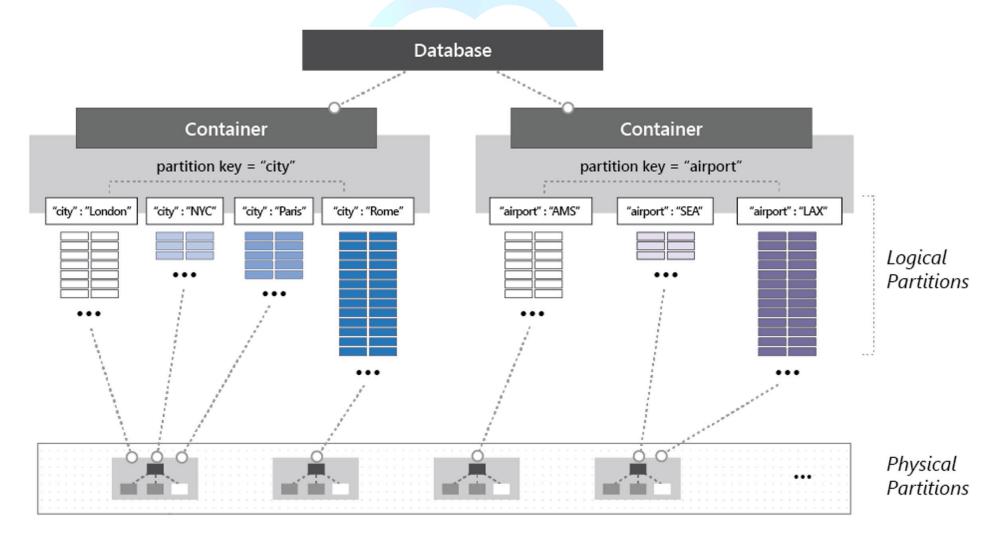


Partitioning



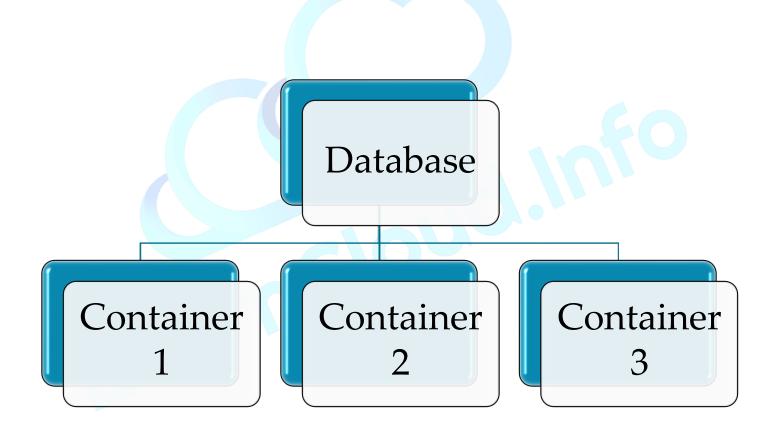


Dedicated vs Shared throughput





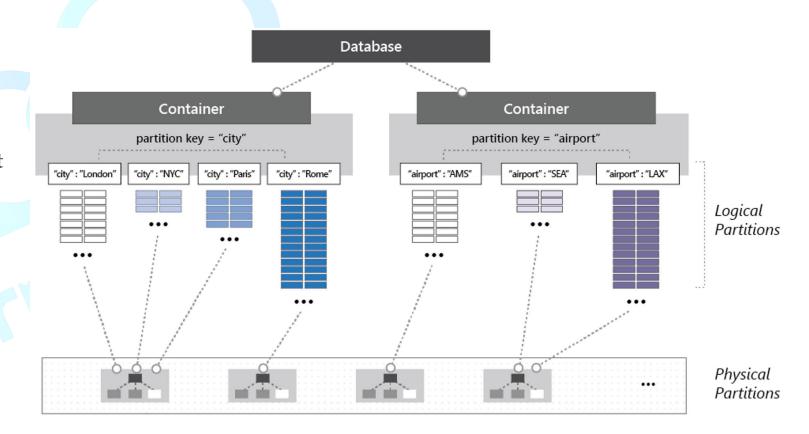
Dedicated vs Shared throughput





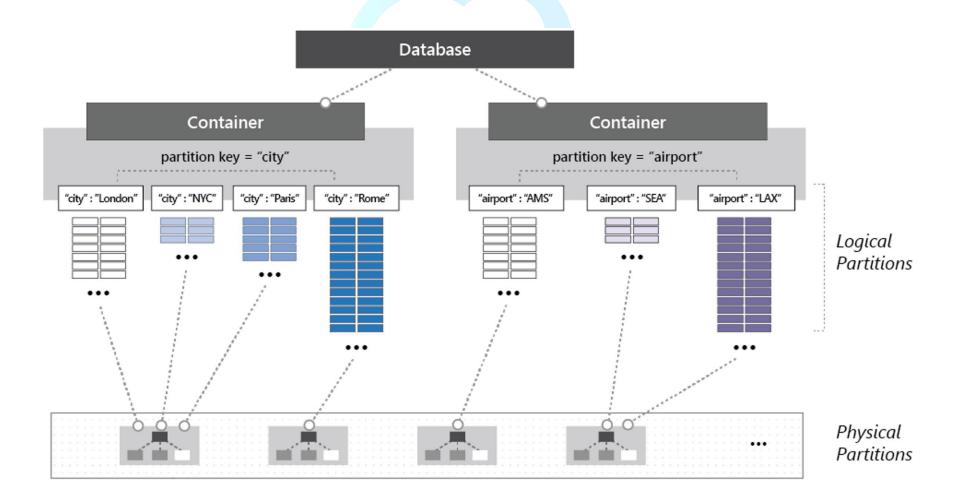
Dedicated vs Shared throughput

- You can set throughput at:
 - Database level Shared throughput
 - Container level Dedicated throughput
 - It is recommend to set throughput at container level.
- Rate-Limited
- Choose at the time of creation





Avoiding hot partition





Avoiding Hot Partitions

Logical Partition 1

(2500 RUs)

Logical Partition 2

(2500 RUs)

Logical Partition 3

(2500 RUs)

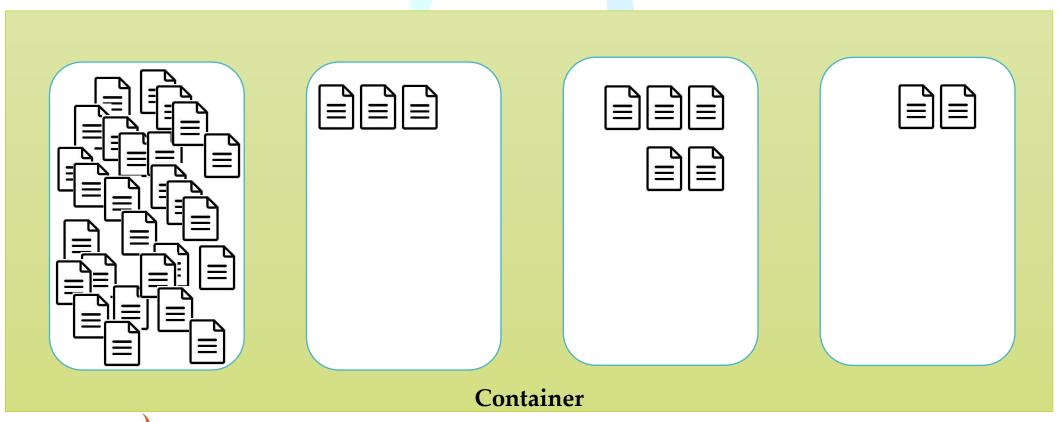
Logical Partition 4

(2500 RUs)

Container (10,000 RUs)



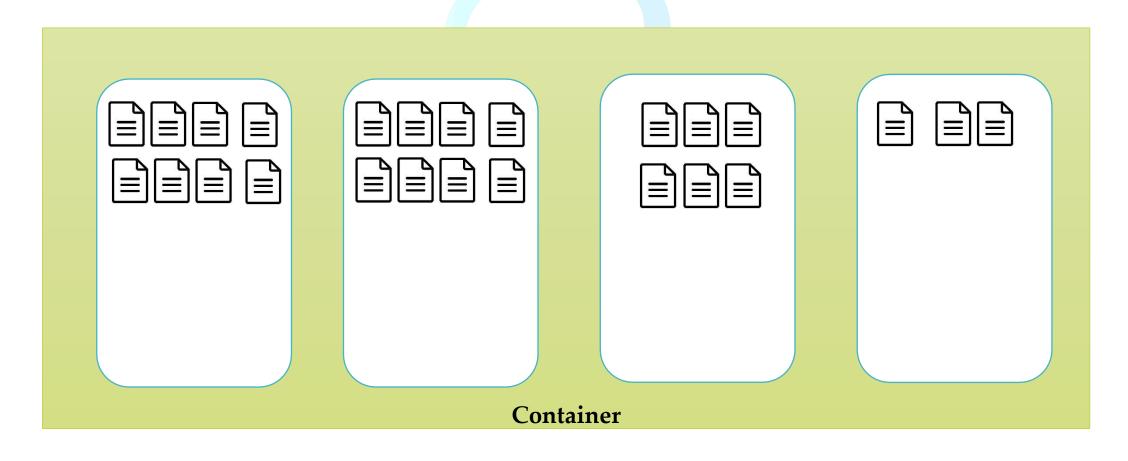
Avoid Hot partitions on storage





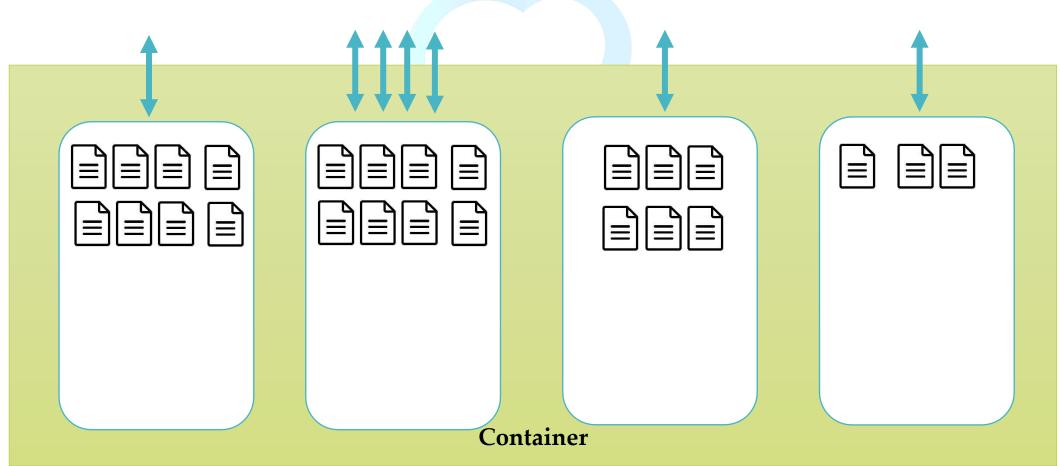


Avoiding Hot Partitions at store





Avoid Hot partitions on throughput

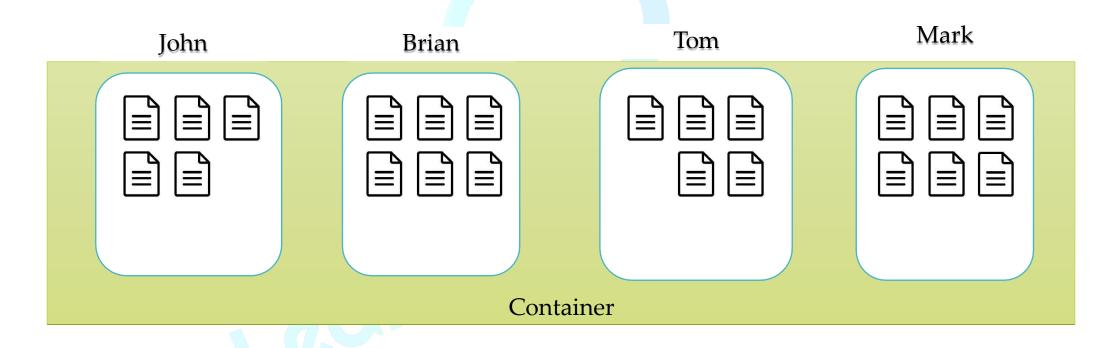




- Partition key Bad choice: Current time
- Partition key Good choices: User ID, Product ID



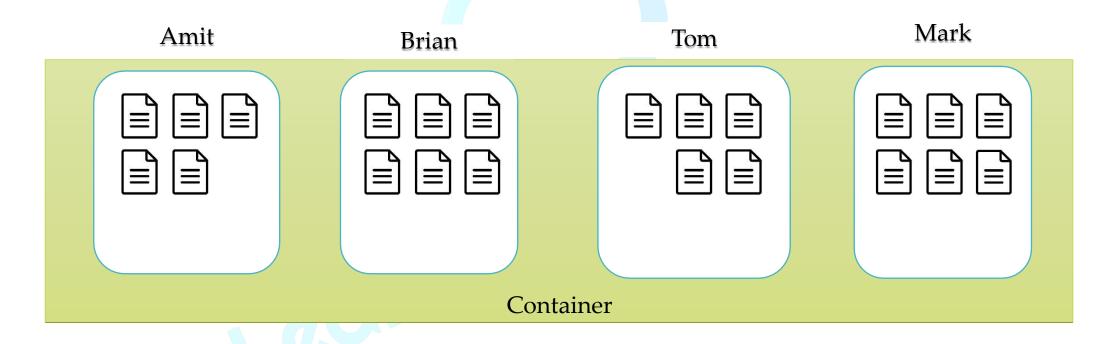
Single partition Query



SELECT * FROM c WHERE c.username = 'Brian'



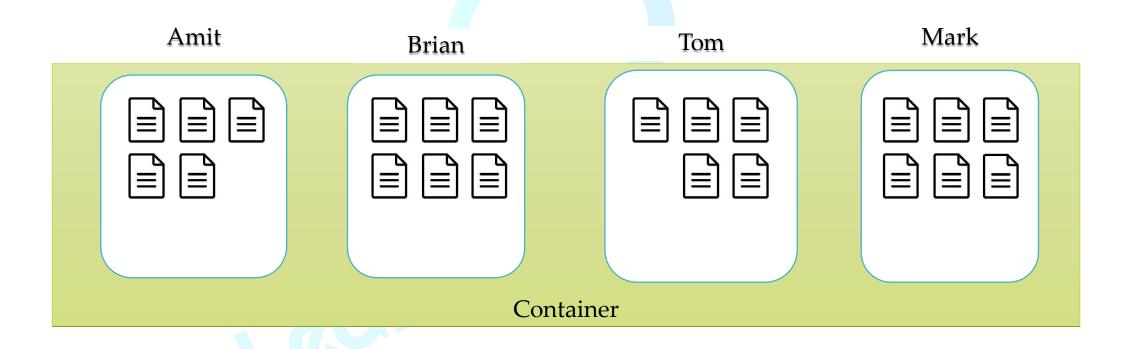
Cross partition Queries (fan out queries)



SELECT * FROM c WHERE c.favoritecolor= 'Blue'



Composite Key



Composite Key: CustomerName-mmddyyyy



Choosing a Partition key

- Evenly distribute storage
 - Make sure you pick your partition key that doesn't result in hot spots within your applications
 - Have a high cardinality
 - Don't be afraid of choosing a partition key that has a large number of values
 - Example User Id & Product Id
- Evenly distribute requests.
 - RUs evenly distribute across all partitions.
 - Review <u>where</u> clause of top queries
- Consider document and partition limit while designing partition key.
 - Max document size 2 MB
 - Max logical partition size 20 GB



Choosing a Partition key

Question: Your organization is planning to use Azure Cosmos DB to store vehicle telemetry data generated from millions of vehicles every second. Which of the following options for your Partition Key will optimize storage distribution?

Answer choices:

- 1. Vehicle model
- 2. Vehicle Identification Number (VIN) which looks like WDDEJ9EB6DA032037



Choosing a Partition key

Question: Your organization is planning to use Azure Cosmos DB to store vehicle telemetry data generated from millions of vehicles every second. Which of the following options for your Partition Key will optimize storage distribution?

Answer choices:



Most auto manufactures only have a couple dozen models. This option is potentially the least granular, will create a

fixed number of logical partitions, and may not distribute data evenly across all physical partitions.



Auto manufacturers have transactions occurring throughout the year. This option will create a more balanced distribution of storage across partition key values.



Automatic Indexing

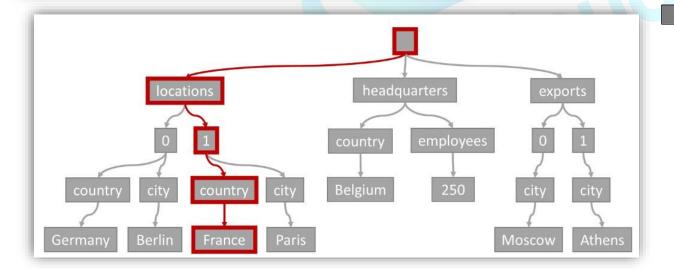
- Index all data without requiring Index management
- Every property of every record automatically index
- Index update synchronously as you create, update or delete items
- Not specific for SQL, but available for all APIs





Automatic Indexing

SELECT location
FROM location IN company.locations
WHERE location.country = 'France'



/locations/0/country: "Germany"
/locations/0/city: "Berlin"
/locations/1/country: "France"
/locations/1/city: "Paris"
/headquarters/country: "Belgium"
/headquarters/employees: 250
/exports/0/city: "Moscow"
/exports/1/city: "Athens"





Time-to-Live

Time to Live (TTL)

- You can set the expiry time for Cosmos DB data items
- Time to live value is configured in seconds.
- The system will automatically delete the expired items based on the TTL value
- Consume only leftover Request units
- Data deletion delay if not enough Request units
 - Though the data deletion is delayed, data is not returned by any queries (by any API) after the TTL has expired.



Global Distribution benefits



Performance

- Ensures high availability within a region
- Across regions, brings data closer to the consumer.

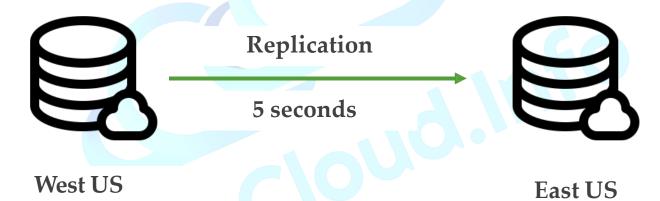


Business continuity

• In the event of major failure or natural disaster



Data consistency

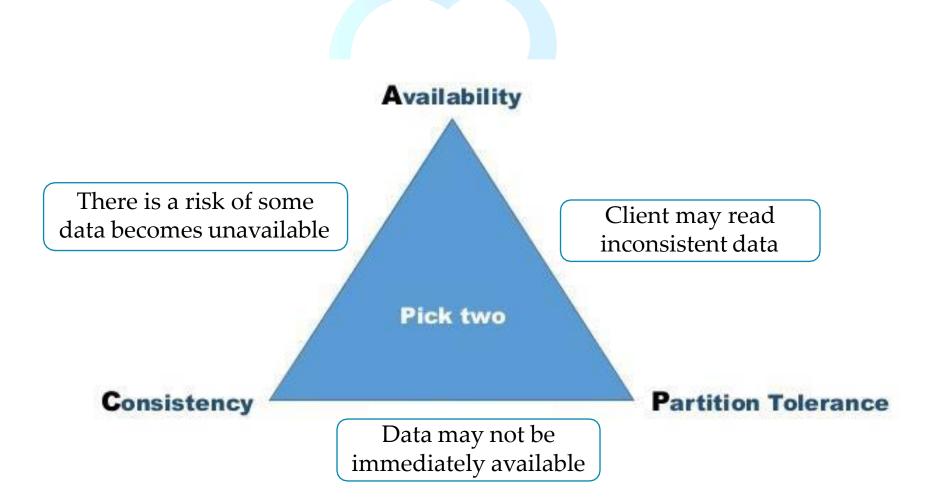


At 10.00 AM Update CreditScore = 750

At 10:00:02 AM Read CreditScore



CAP Theorem





Five consistency Levels



Strong: No dirty reads, high latency, cost highest, closest to RDBMS

Bounded staleness: Dirty reads possible, bounded by time and updates

Session: No dirty reads for writers (within same session), dirty read possible for other users

Consistency prefix: Dirty reads possible but sequence maintain, reads never see out-of-order writes

Eventual: No guaranteed order, but eventually everything gets in order



Setting the consistency level

Set default for entire account

Can be changed any time

Override at request level

Any request can weaken the default consistency level

Strong Consistency Higher latency Lower availability



Bounded Staleness

Session

Consistent Prefix

Eventual

Weaker Consistency Lower latency Higher availability



Azure CLI

```
az cosmosdb create -- name
                   --resource-group
                   [--capabilities]
                   [--default-consistency-level {BoundedStaleness, ConsistentPrefix, Eventual, Session, Strong}]
                   [--enable-automatic-failover {false, true}]
                   [--enable-multiple-write-locations {false, true}]
                   [--enable-virtual-network {false, true}]
                   [--ip-range-filter]
                   [--kind {GlobalDocumentDB, MongoDB, Parse}]
                   [--locations]
                   [--max-interval]
                   [--max-staleness-prefix]
                   [--subscription]
                   [--tags]
                   [--virtual-network-rules]
```



Azure CLI - Create Database example

```
reate a SQL API database and container
# Generate a unique 10 character alphanumeric string to ensure unique resource names
uniqueId=$(env LC_CTYPE=C tr -dc 'a-z0-9' < /dev/urandom | fold -w 10 | head -n 1)
# Variables for SQL API resources
resourceGroupName="Group-$uniqueId"
location='westus2'
accountName="cosmos-$uniqueId" #needs to be lower case
databaseName='database1'
containerName='container1'
# Create a resource group
az group create -n $resourceGroupName -1 $location
# Create a Cosmos account for SQL API
az cosmosdb create \
    -n $accountName \
    -g $resourceGroupName \
    --default-consistency-level Eventual \
    --locations regionName='West US 2' failoverPriority=0 isZoneRedundant=False \
    --locations regionName='East US 2' failoverPriority=1 isZoneRedundant=False
# Create a SQL API database
az cosmosdb sql database create \
    -a $accountName \
    -g $resourceGroupName \
    -n $databaseName
```



Azure CLI - Example question

You are a data engineer for your company. You use the following Azure CLI commands to create an Azure Cosmos DB account. You plan to use this account to store sales data.

az cosmosdb create --resource-group 'sales-rg' --name 'sales' --kind GlobalDocumentDB \

- --locations regionName="South Central US" failoverPriority=0 \
- --locations regionName="North Central US" failoverPriority=1 \
- --default-consistency-level "Strong" --enable-multiple-write-locations true

You need to answer questions regarding sales data queries and updates.

For each of the following statements, select Yes if the statement is true. Otherwise, select No.

Statement



You can query data by using Gremlin API.



A client can see partial writes of a sales data record by default.



A client can set the consistency level to Eventual Consistency at connection time.



A client can set a different consistency level during each request to sales data.

